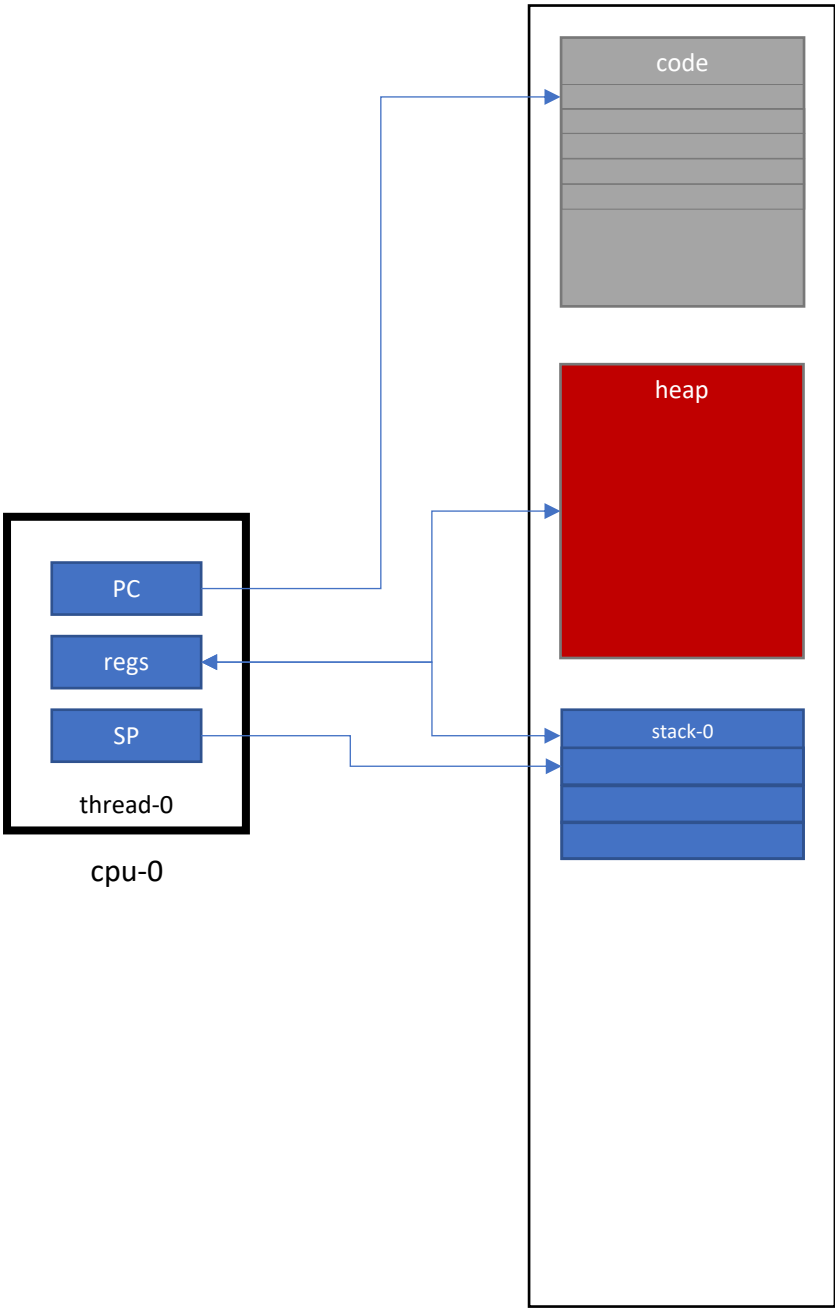
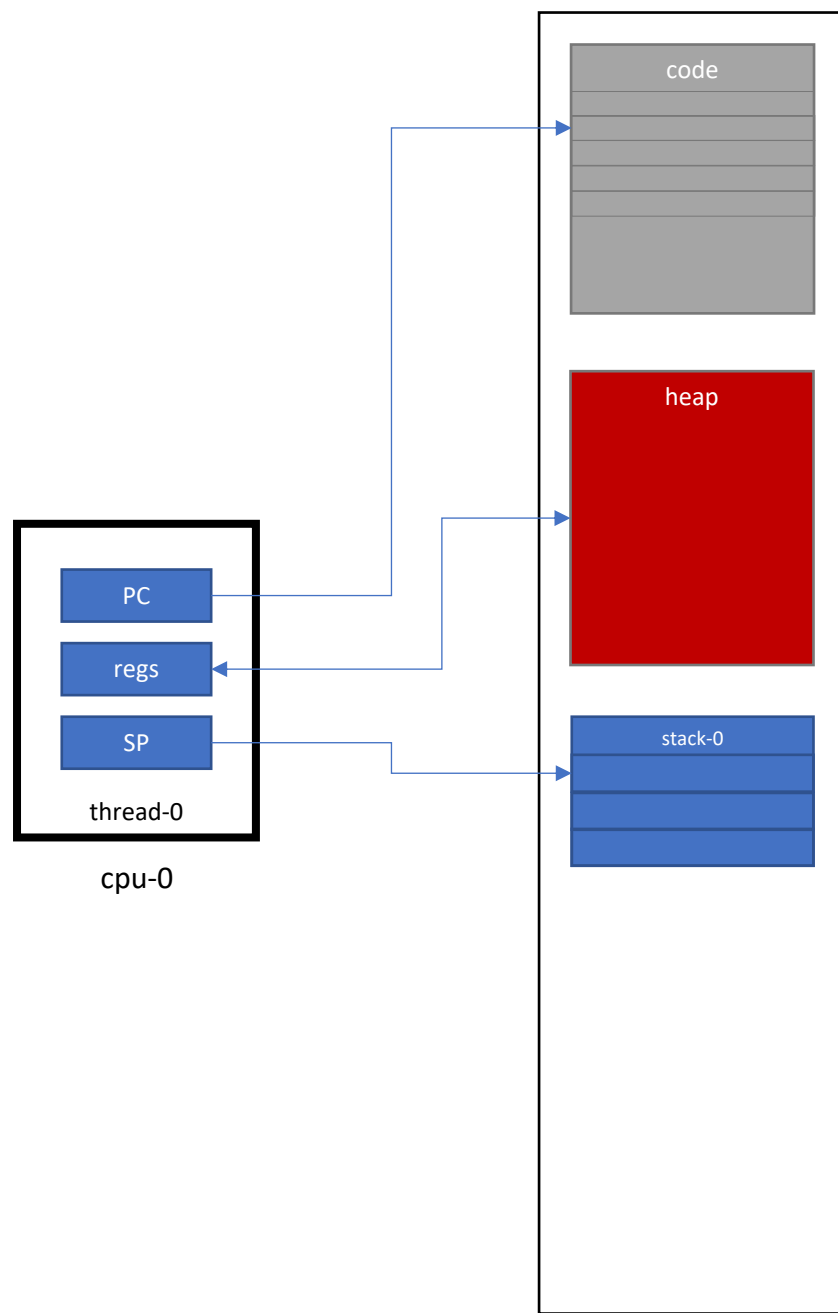


Threads

Concurrent Programming





Same **code** block

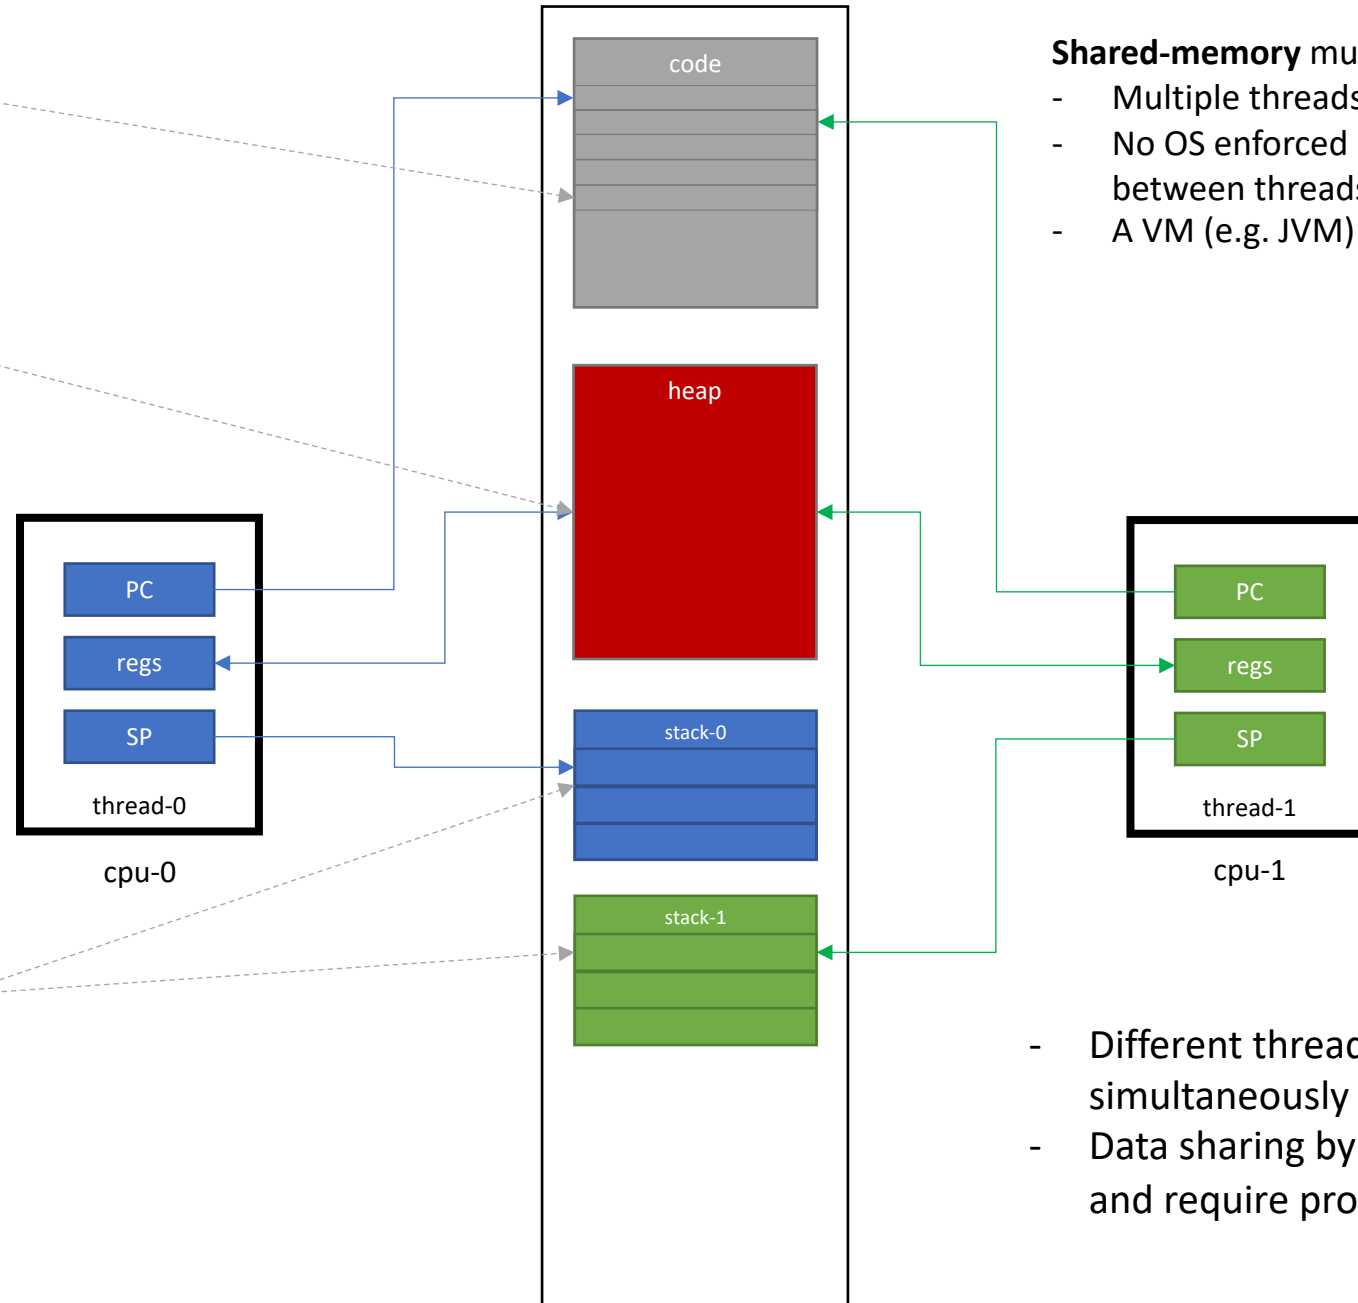
Same **heap** block

Distinct **stack** blocks

stack contains:

- return addresses
- local variables
- function parameters

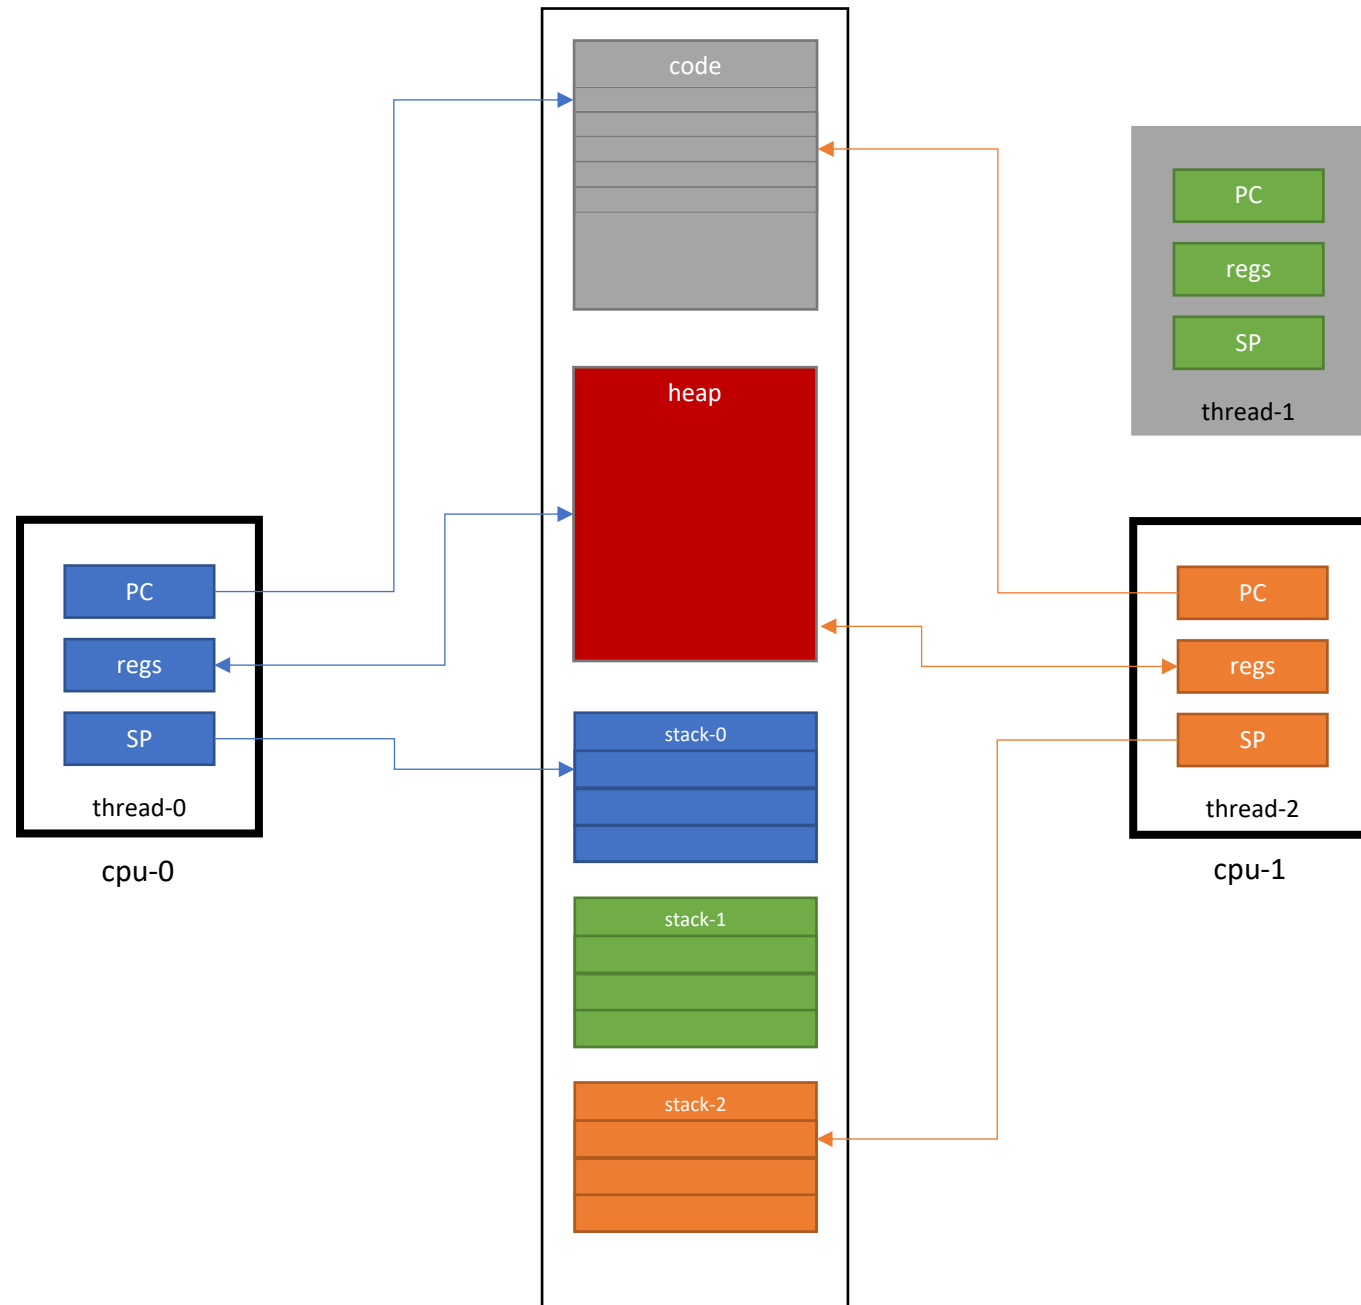
specific to each thread



Shared-memory multi-threading:

- Multiple threads share the same code and heap
- No OS enforced memory protection mechanism between threads
- A VM (e.g. JVM) can impose some restrictions

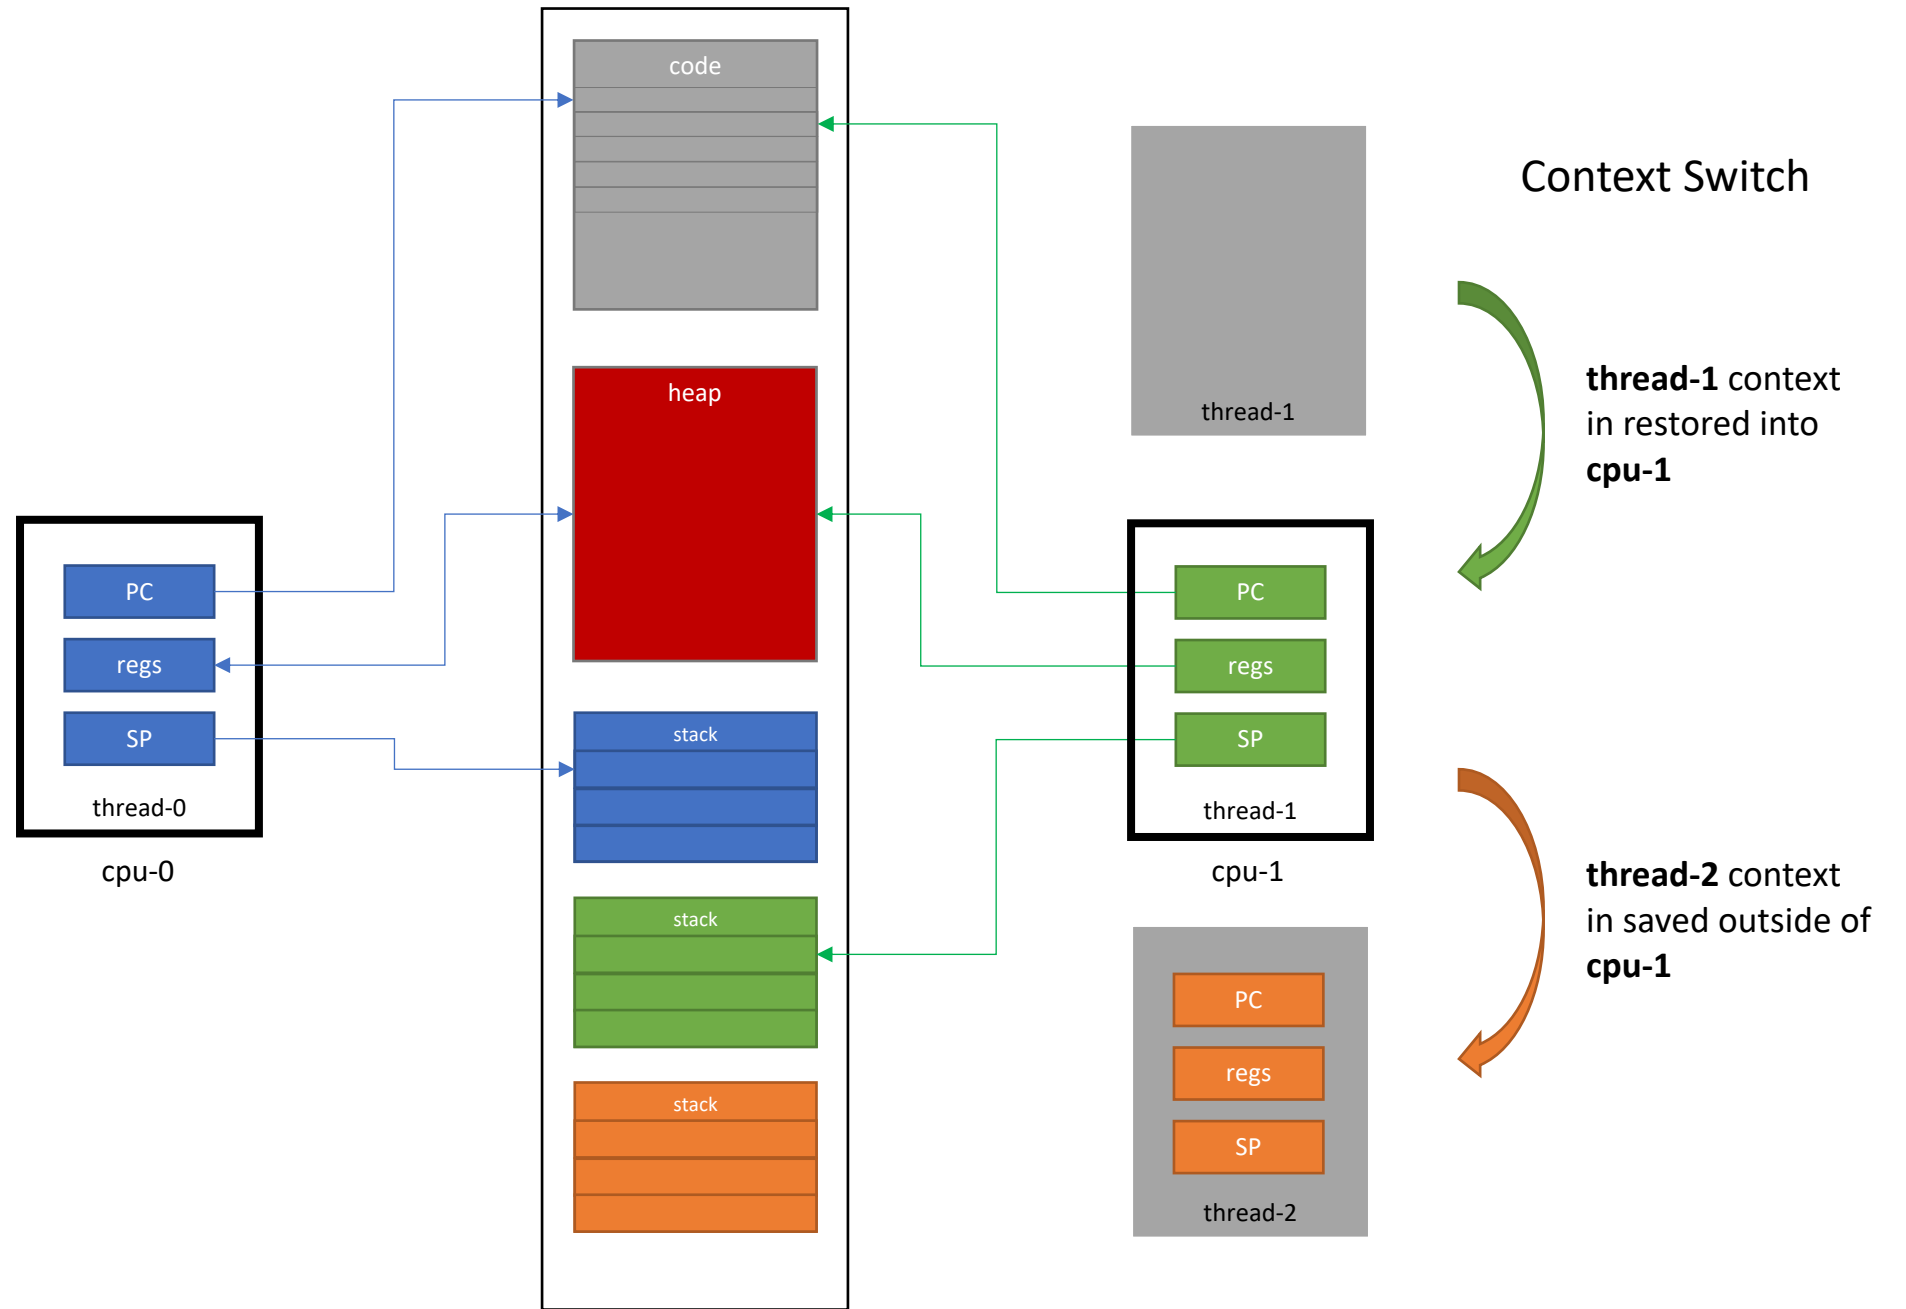
- Different threads can be *running* simultaneously on different CPUs
- Data sharing by different threads is *hazardous* and require proper *synchronization* techniques

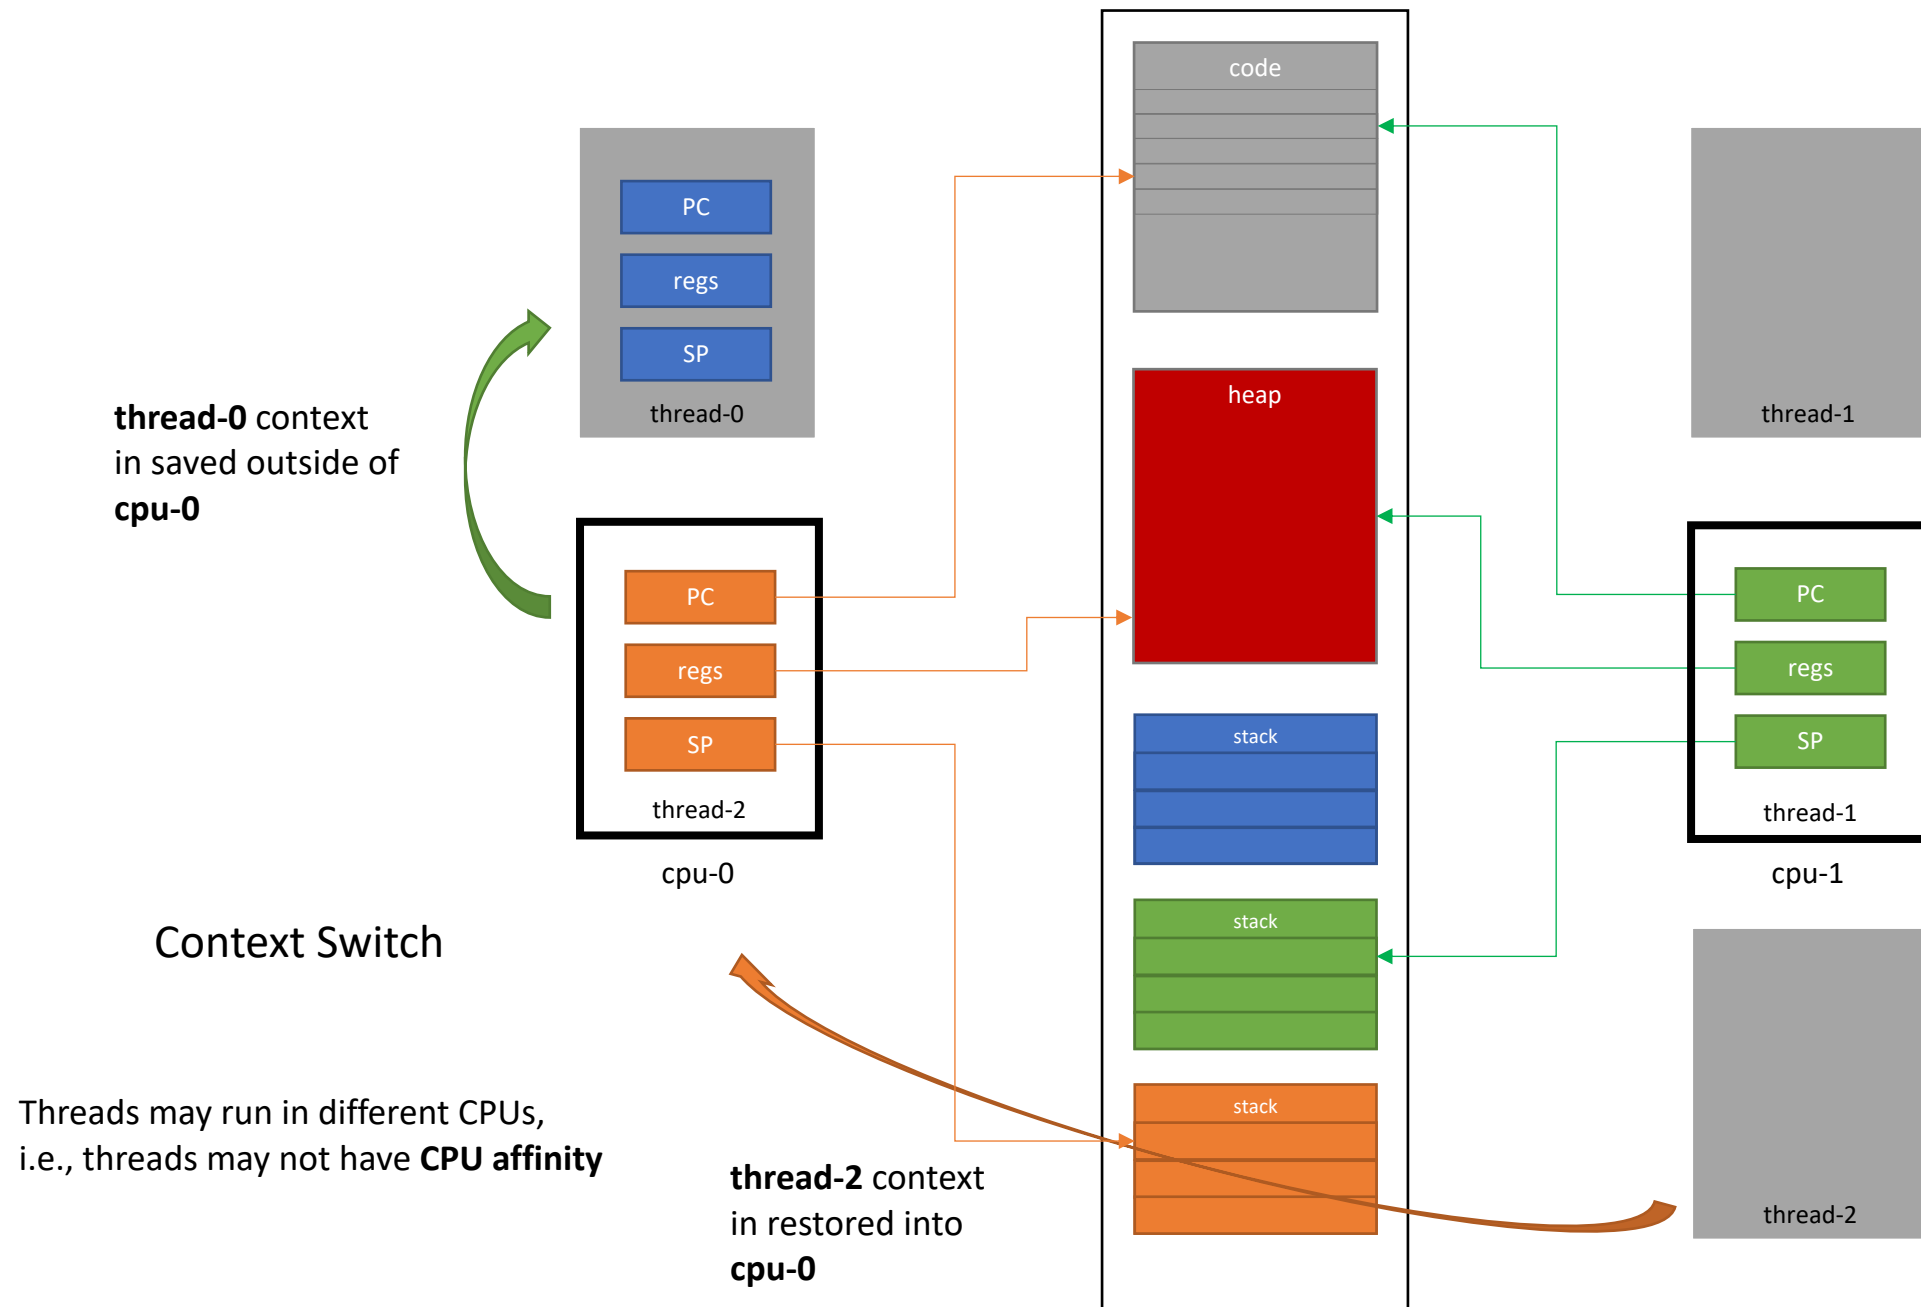


There can be more threads than processors

- **thread-1 context** is saved outside of **cpu-1**
- **thread-2** starts executing on **cpu-1**

thread-2 has its own stack, separate from **thread-1** stack



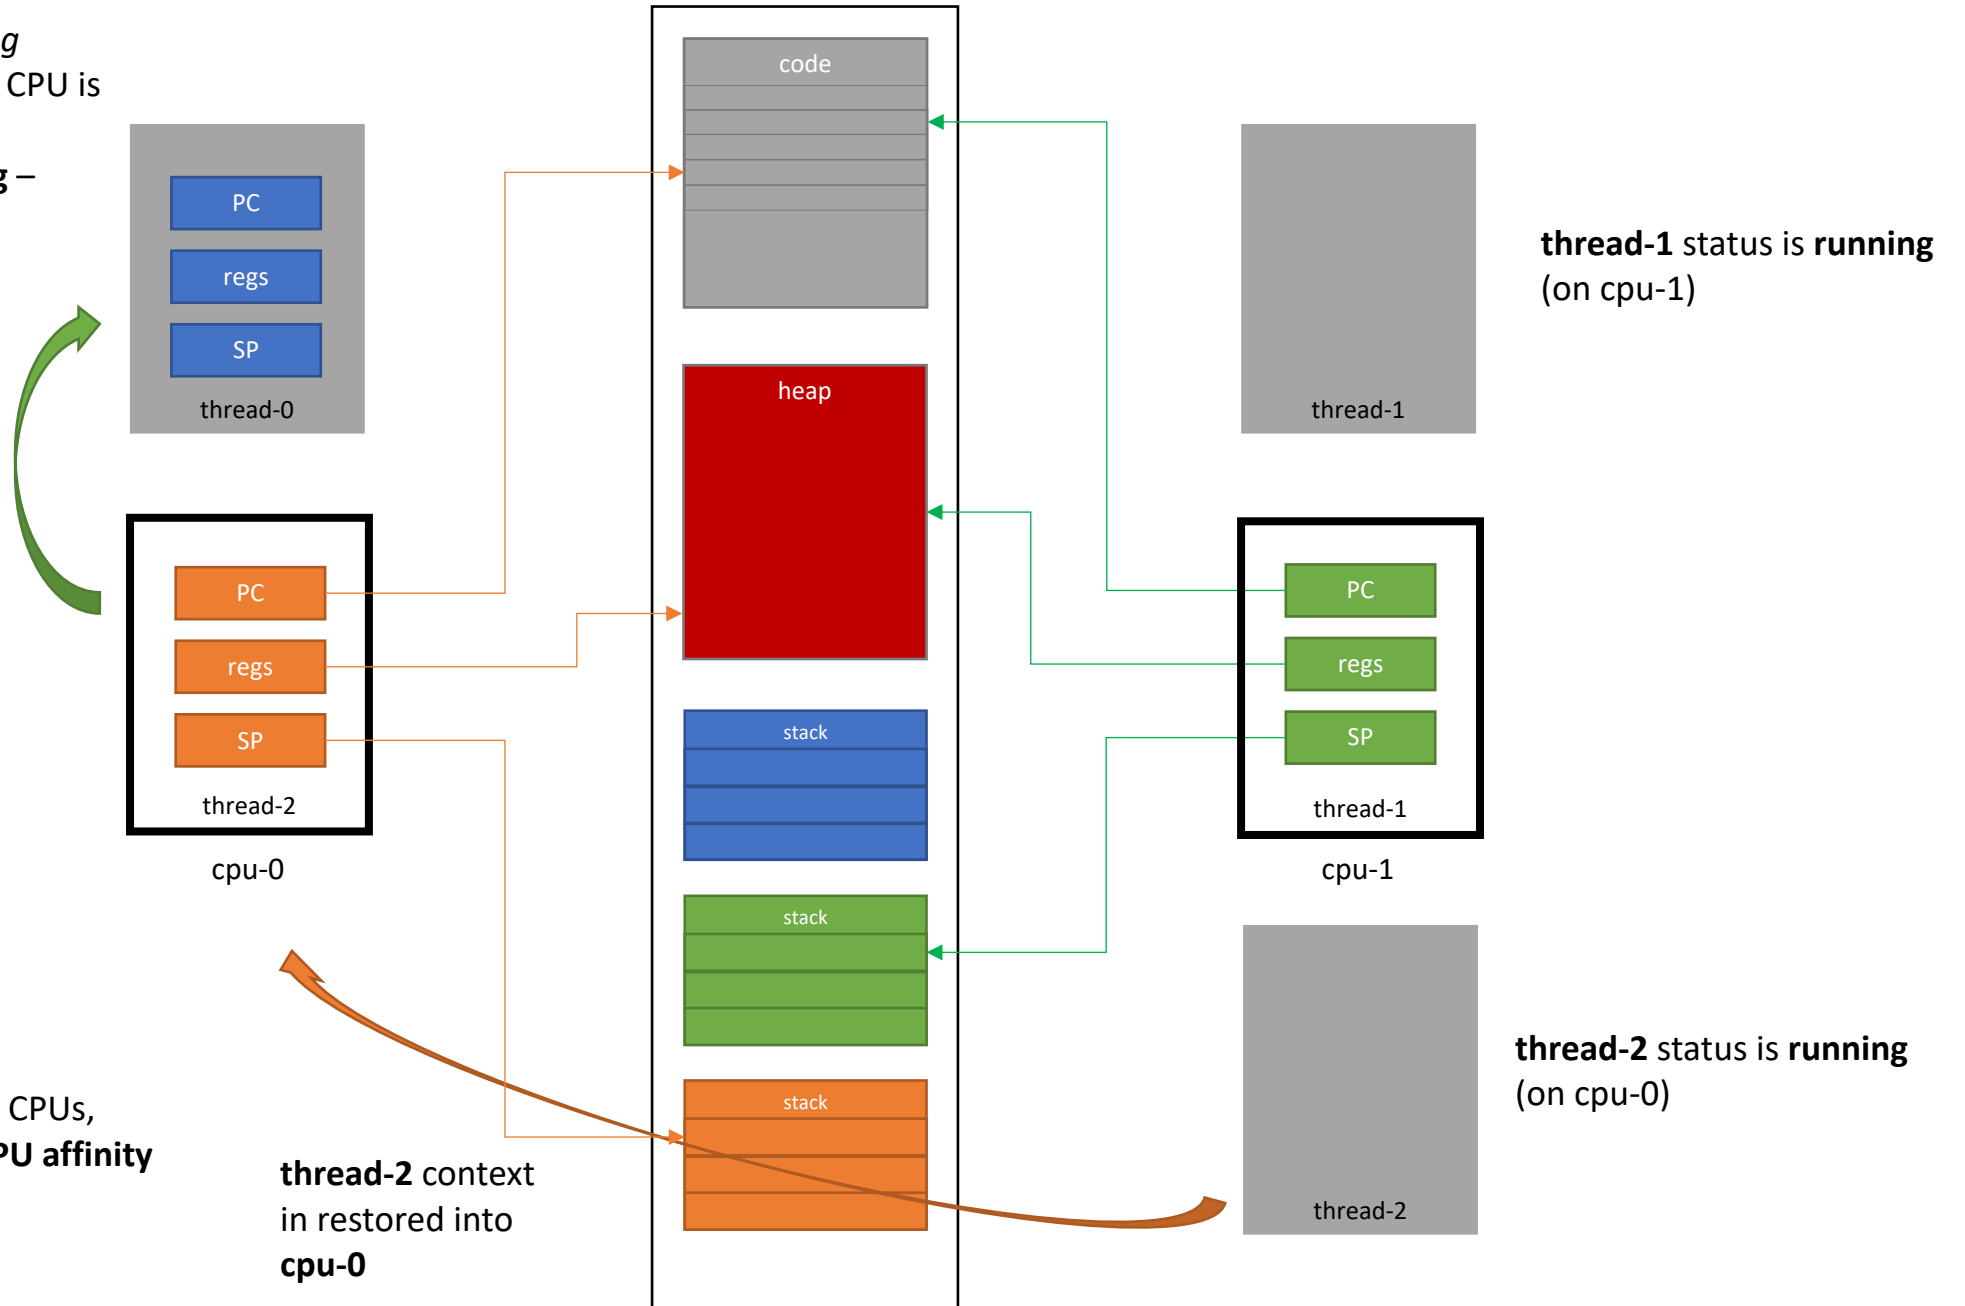


thread-0 status is *not running*
- **ready** – able to run when a CPU is available
- **not-ready/blocked/waiting** – not able to run, waiting for *something* to happen

thread-0 context is saved outside of **cpu-0**

Context Switch

Threads may run in different CPUs, i.e., threads may not have **CPU affinity**



Scheduling

- Conceptually, there is a list with all the **ready** threads
- Scheduling is the process
 - Deciding which **running** threads should transition to **ready** (i.e. *lose* the CPU)
 - Deciding which **ready** threads should transition to **running** (i.e. *gain* the CPU)
 - Performing the associated context switches
- When is scheduling performed?
 - When a thread calls the scheduling code (e.g. indirectly via a function call)
 - When an interrupt (system timer) occurs and calls the scheduler
 - This means that code running on a thread cannot control the points where the scheduling occurs and the thread is switched out of the CPU.
 - I.e. a thread can be switched out at any assembly instruction boundary.

Scheduling

- Scheduling is based on
 - Thread **status**: only **running** and **ready** threads are eligible to go/keep **running**
 - Thread **priority**: a way of sort to prefer some threads in relation to others
 - Because some threads may be more “important” than others
 - Thread **execution time**: how long has a thread been in the running state
 - To do time multiplexing between ready and running threads.

OS threads vs. VM threads

- VM-based programming environments, such as the JVM and .NET, provide a way to create threads associated with managed code.
- Typically, these *managed threads* are implemented using OS threads
 - 1-1 model: one managed thread uses one OS thread.
 - Scheduling is performed by the OS, which controls the OS threads.
- In the JVM, this model may change with Project Loom - <https://openjdk.java.net/projects/loom/>

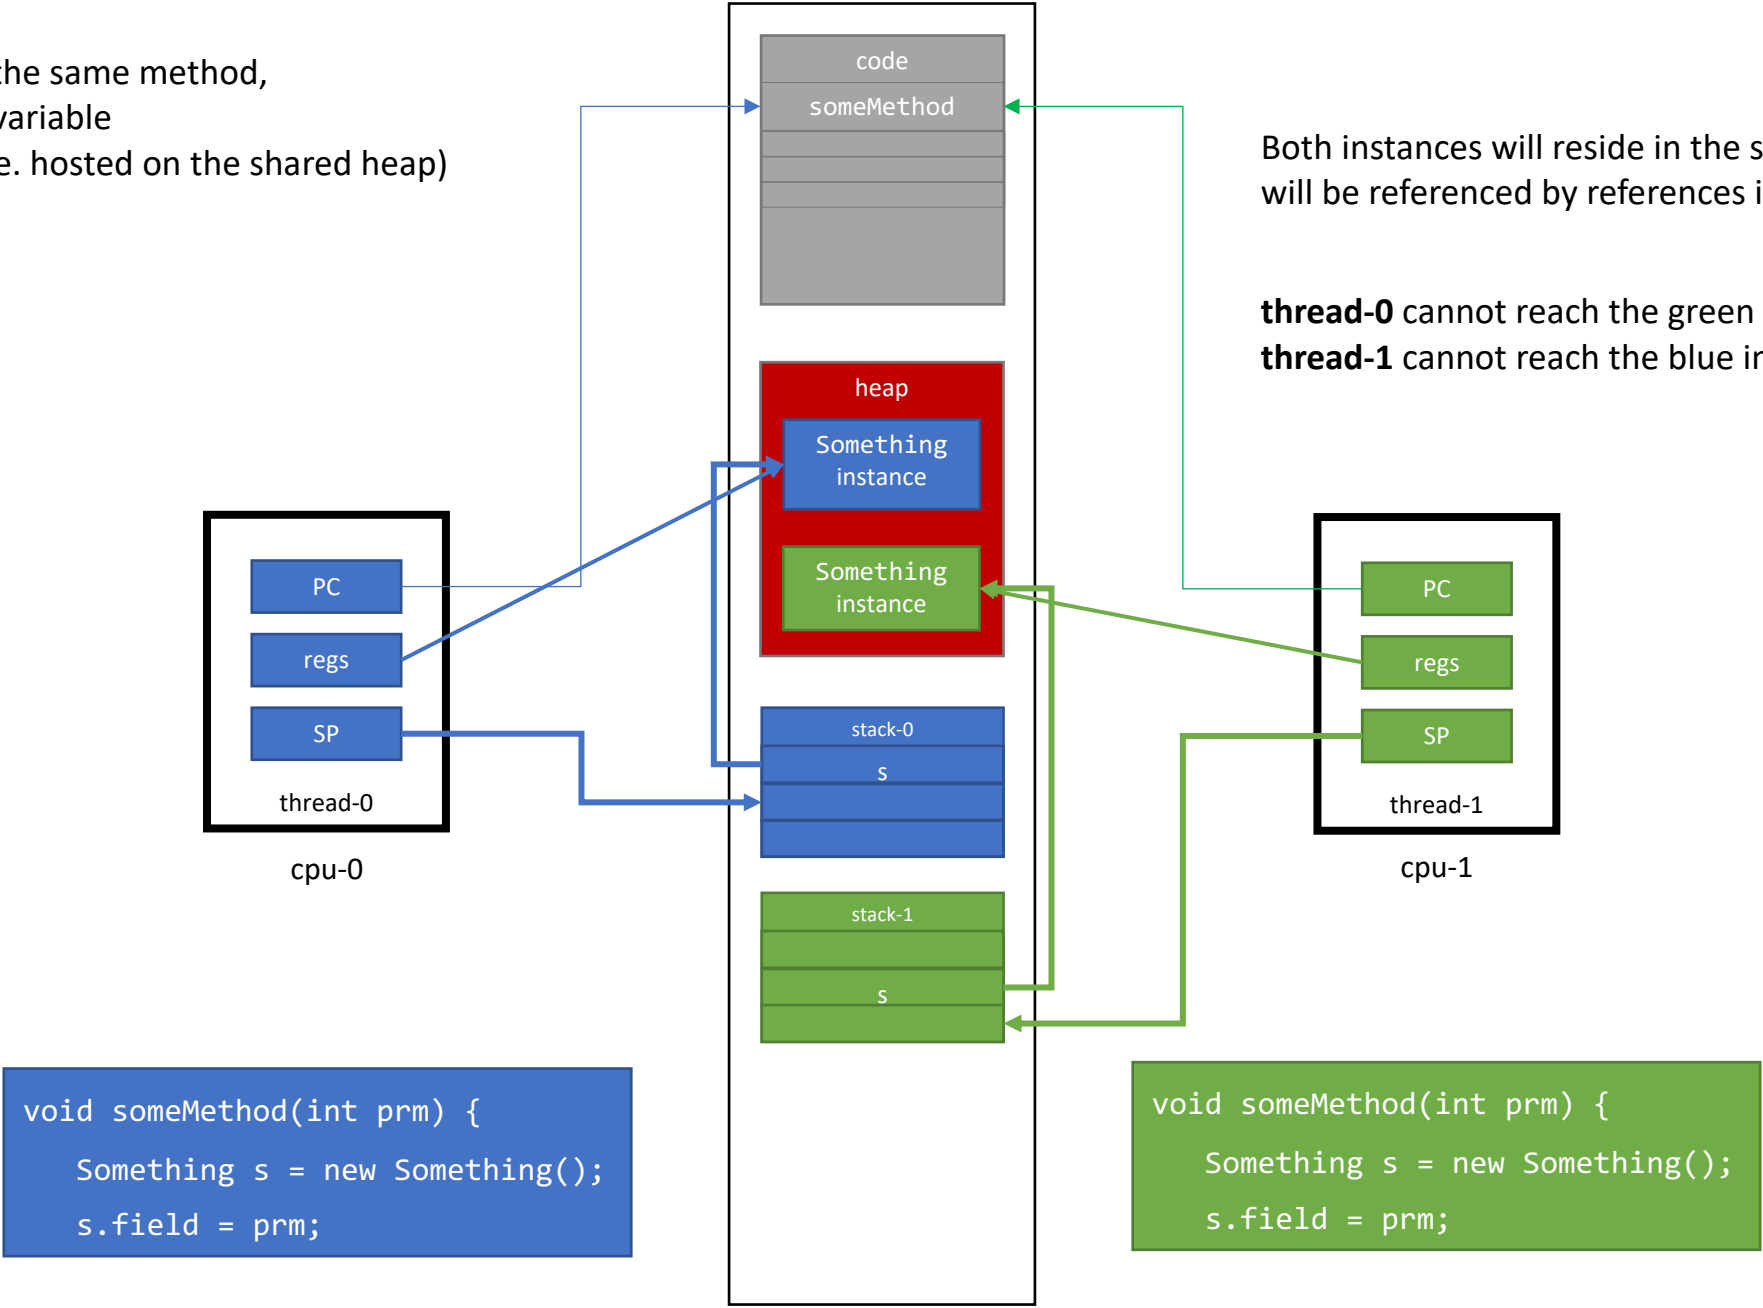
Data sharing

- A shared-memory multi-threading model doesn't mean all data is effectively shared.
- Identifying and avoid unnecessary sharing is a very important skill
- Let's look at some examples...

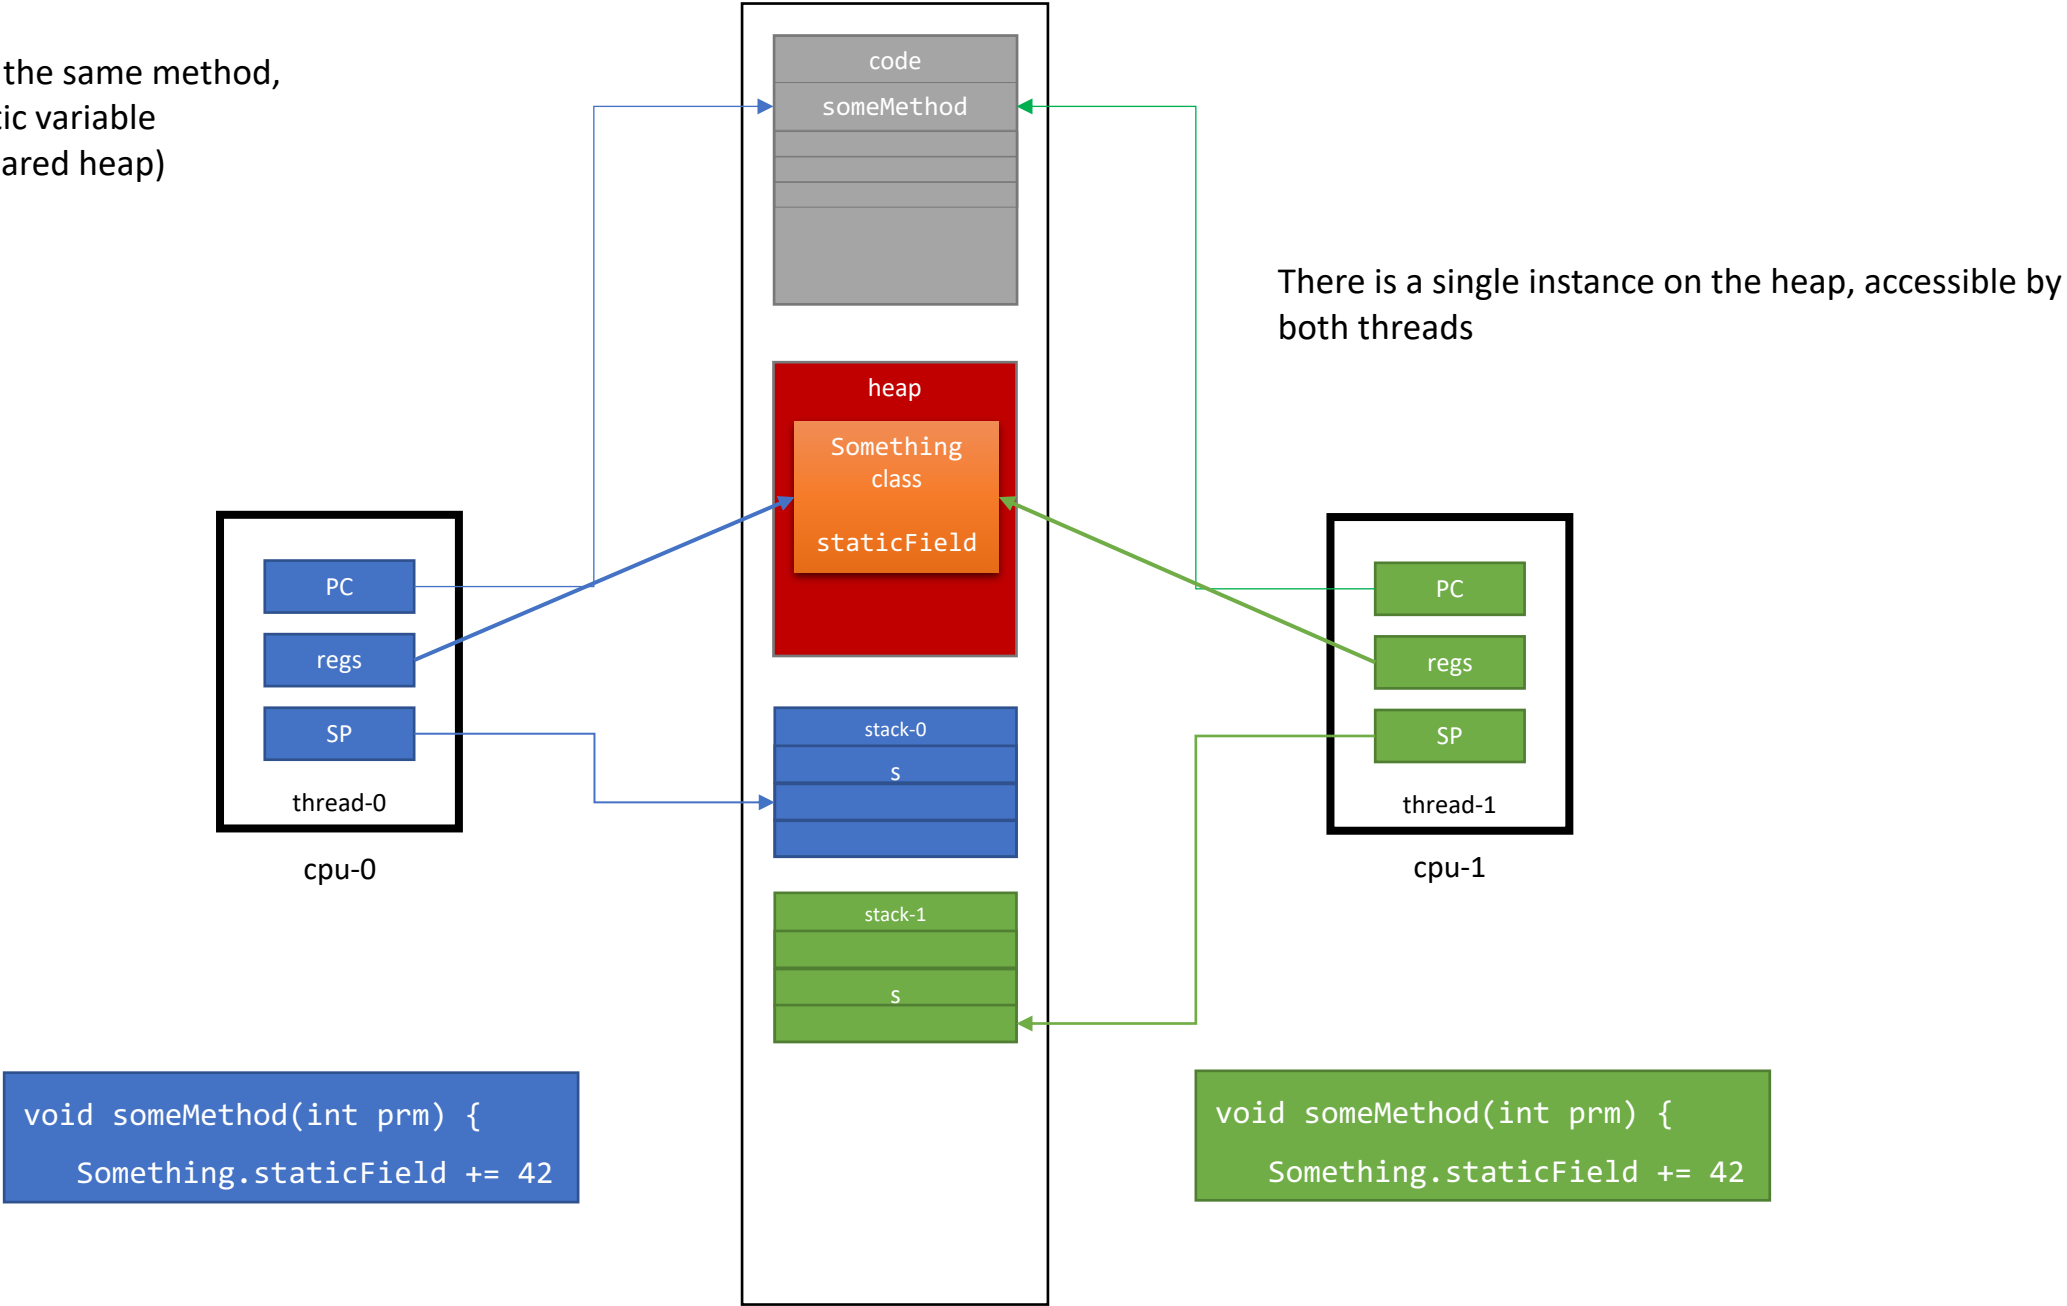
Two threads running the same method,
which creates a local variable
of a reference type (i.e. hosted on the shared heap)

Both instances will reside in the shared heap, however
will be referenced by references in distinct stacks.

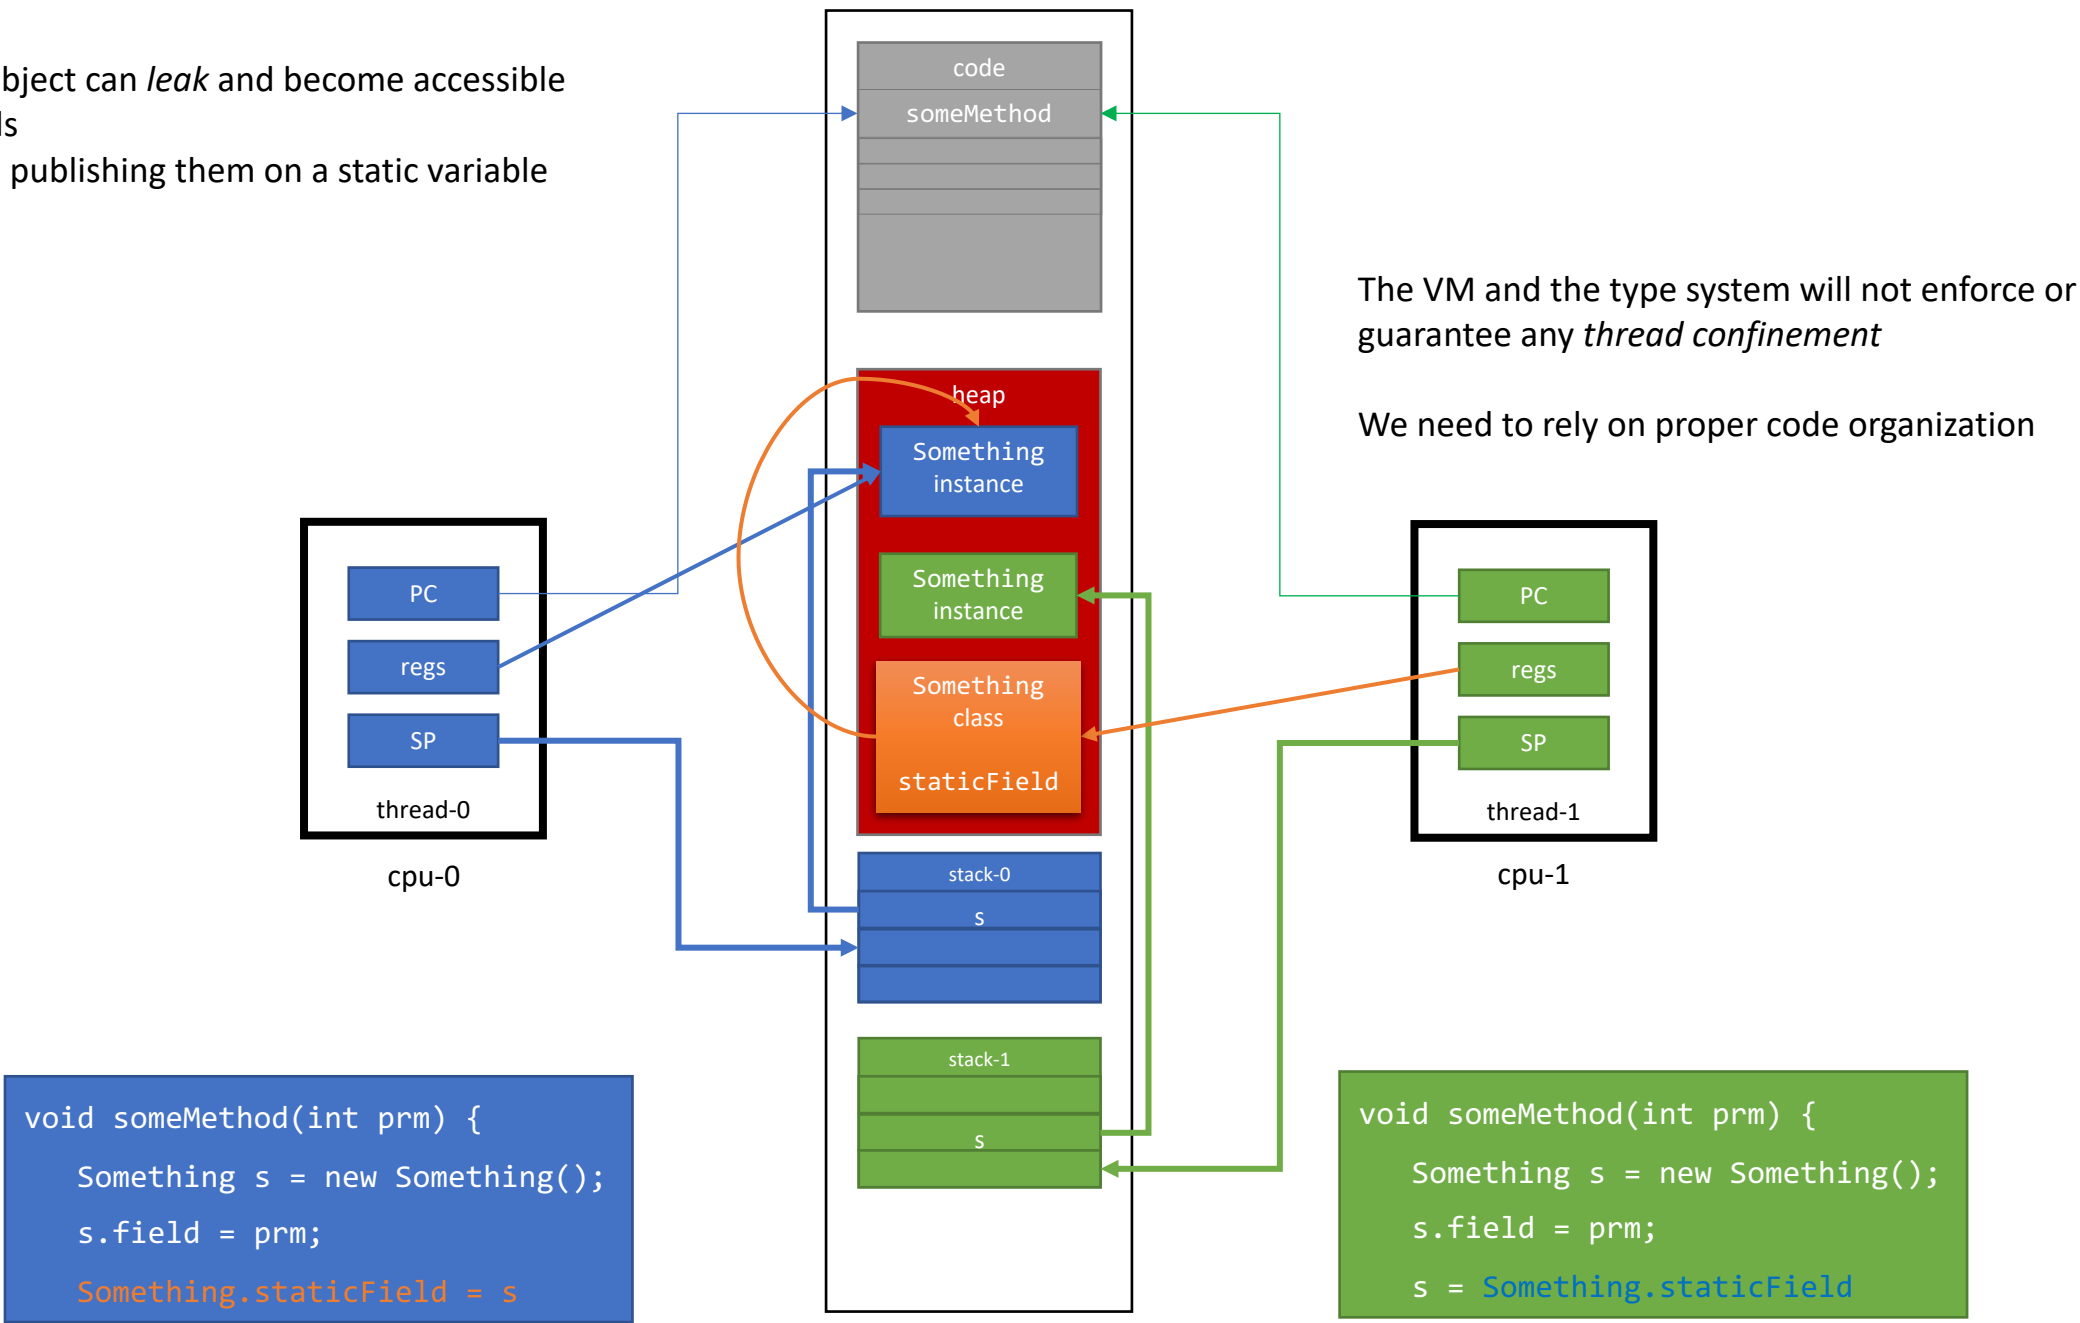
thread-0 cannot reach the green instance
thread-1 cannot reach the blue instance



Two threads running the same method,
which accesses a static variable
(i.e. hosted on the shared heap)



References to object can *leak* and become accessible by other threads
E.g. by a thread publishing them on a static variable

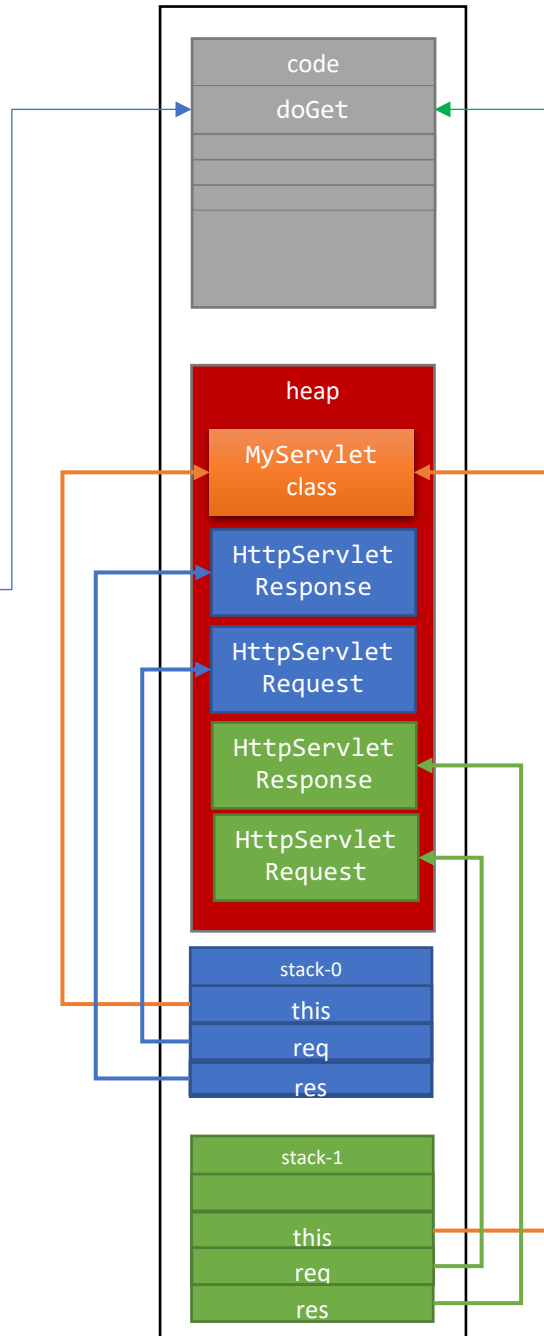



Recalling the LS project...

Example:


- A single **MyServlet** instance, used to handle multiple simultaneous HTTP requests

```
class MyServlet extends HttpServlet {  
    private Something aField;  
    @Override  
    void doGet(  
        HttpServletRequest req,  
        HttpServletResponse res) {  
    }  
}
```



Each thread will access

- A shared **MyServlet** instance fields
- Request specific (and therefore thread specific) request and response objects




```
class MyServlet extends HttpServlet {  
    private Something aField;  
    @Override  
    void doGet(  
        HttpServletRequest req,  
        HttpServletResponse res) {  
    }  
}
```

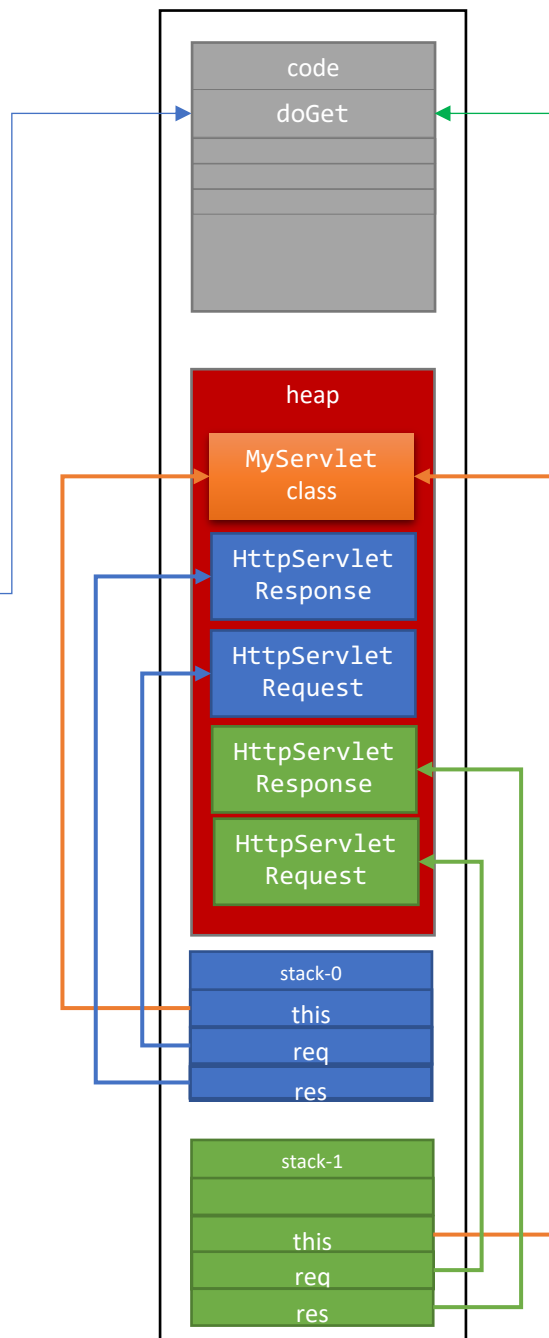
Example:

- A single **MyServlet** instance, used to handle multiple simultaneous HTTP requests

The type system
or the VM
will **not** warn us of this




```
class MyServlet extends HttpServlet {  
    private HttpServletRequest m_req;  
    @Override  
    void doGet(  
        HttpServletRequest req,  
        HttpServletResponse res) {  
        m_req = req;  
    }  
}
```



Each thread will access

- A shared **MyServlet** instance fields
- Request specific (and therefore thread specific) request and response objects



```
class MyServlet extends HttpServlet {  
    private HttpServletRequest m_req;  
    @Override  
    void doGet(  
        HttpServletRequest req,  
        HttpServletResponse res) {  
        m_req = req;  
    }  
}
```

Three types of data

- Mutable thread-bound data
 - E.g. the servlet request and response objects
 - No thread sharing issues (because no sharing)
- Immutable shared data
 - E.g. the servlet, the router, the handlers
 - No thread sharing issues (if *some* requirements are fulfilled)
- Mutable shared data
 - E.g. the data source
 - Prone to concurrency hazards
 - Proper synchronization is required